

Inference Latency: Cascade vs Direct

